

ARCMiS

Automated Rust migration for legacy C, Java, and Go codebases

ARCMiS team

Coverage sweep status: 2026-10-04

Contents

1. Introduction and problem description
2. Methodology
3. Inspirations
4. Benchmarks
5. References
6. Thank you

Introduction and problem description

The migration problem

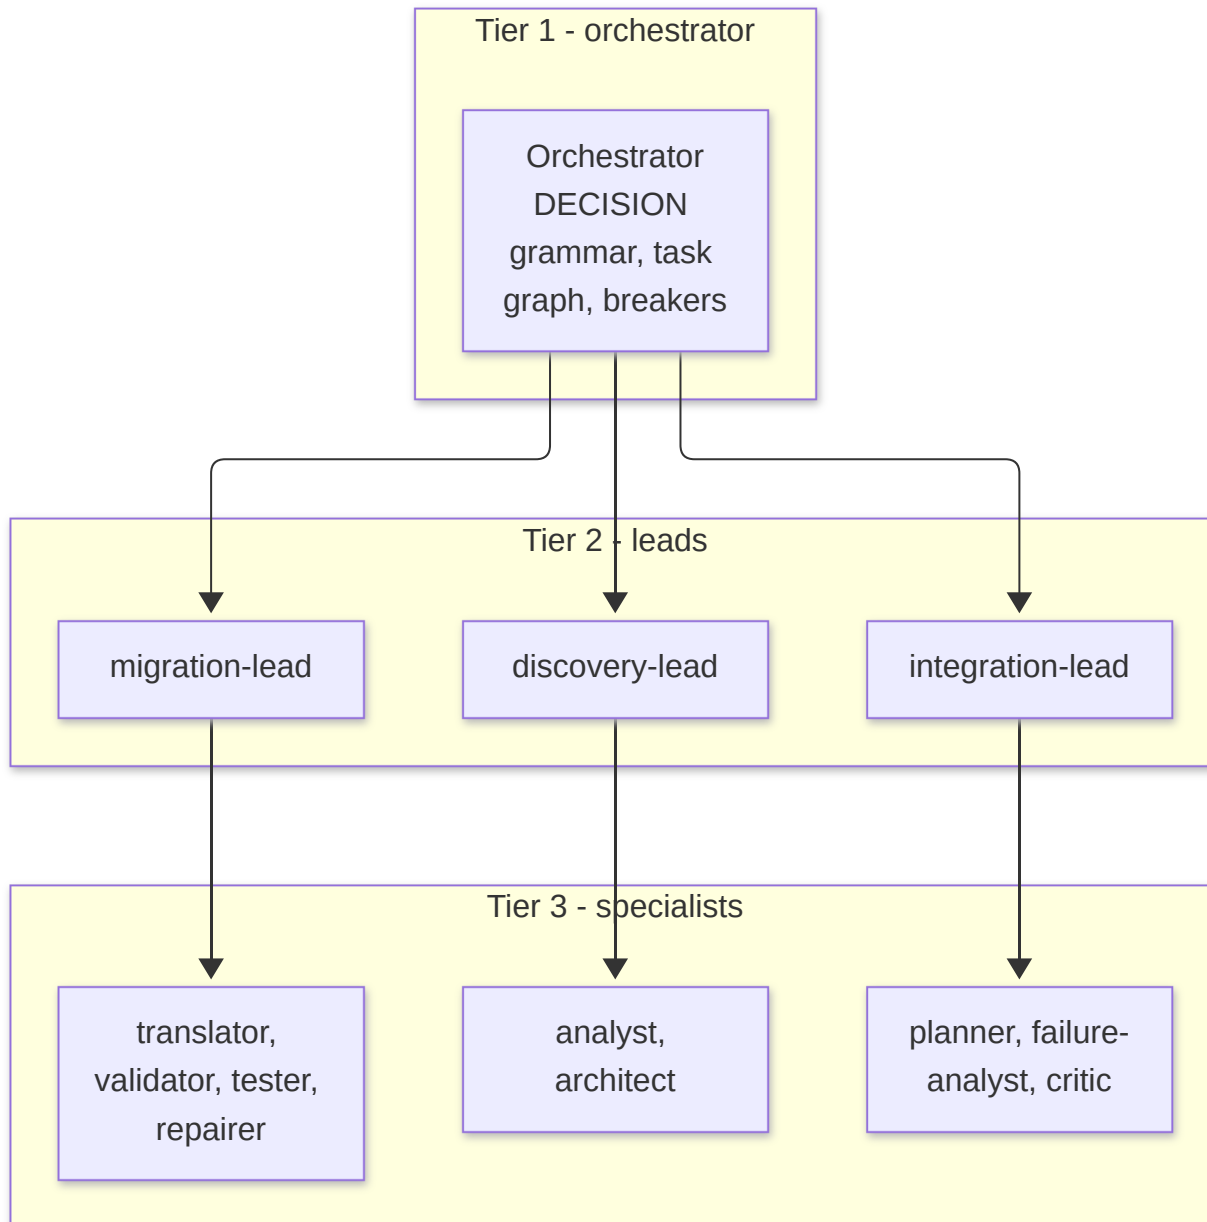
- Real codebases still run on legacy C, Java, and Go.
- Safe, memory-checked Rust is the target language for new work.
- Hand migration is slow. Each project needs source reading, planning, translation, and test validation.
- A local 27B model can drive this work. A single LLM call cannot.
- The ARCMiS answer: a harness of cooperating agents, a fixed problem set, and toolchain-only scoring.

What ARCMiS is

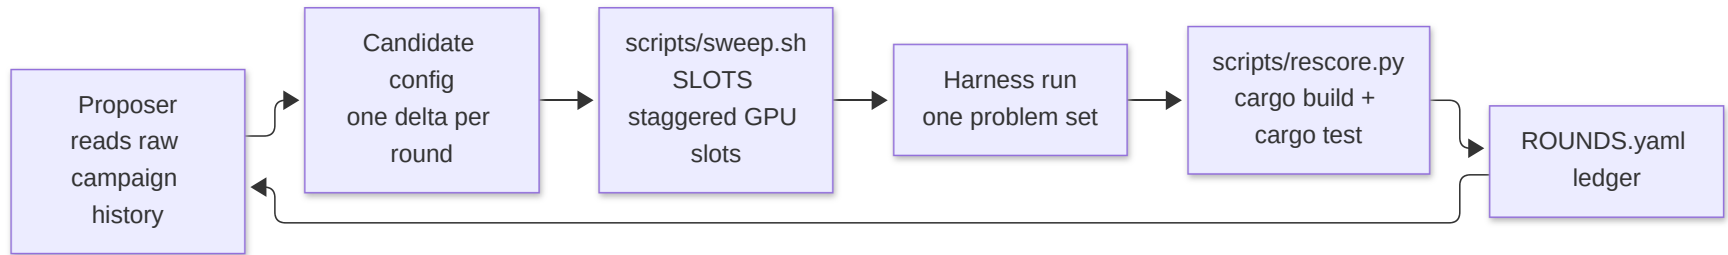
- A code-migration harness over the rig agent framework in Rust.
- Hierarchical multi-agent orchestration (ADR 0022, ADR 0026).
- One candidate config per problem set, one set at a time, LoC ascending.
- Every verdict comes from the toolchain: `cargo build` and `cargo test` through `scripts/rescore.py` .
No model self-grading.
- Campaign method: the Meta-Harness loop (arXiv:2603.28052) proposes, runs, and scores harness changes against a fixed problem set.

Methodology

The MAS hierarchy



The coverage sweep pipeline



- `scripts/sweep.sh` runs waves of `SLOTS` (default 2) staggered GPU slots, plus a CPU proposer and scorer pipeline. One candidate per set.
- Pairs: one crust rust-variant plus one non-crust set. Per-problem deltas live in `scripts/overrides/`.
- Coverage rule: untouched projects only. A set counts once across retries.
- Verdict from `scripts/rescore.py` : toolchain-only scoring.

Inspirations

Orchestration patterns: the 2-axis taxonomy

Source: "LLM-Based Multi-Agent Orchestration: A Survey of Frameworks, Communication Protocols, and Emerging Patterns" (the Q2 survey).

Axes: centralized / decentralized / hierarchy x static / dynamic-adaptive.

ARCMiS position: hierarchy x dynamic-adaptive.

- Hierarchy: tier-2 leads over tier-3 specialists. Verified in `ARCMiS/lib/orchestrator/src/hierarchy.rs` (tier-2 team resolution, config-owned under `mas.hierarchy`) and `lead.rs` (lead inner loop). Recorded in ADR 0022 and ADR 0026.
- Dynamic: the task graph re-scores after each delegation, the router picks roles on intermediate state, and the orchestrator can spawn a failure-analyst sub-chain mid-task.
- Static agents: fixed registry, fixed preambles, fixed tool allowlists. Selection is dynamic, not the agents.

Modernization workflows

ReCodeAgent (arXiv:2604.07341, ASE 2026, vendored at `assets/ReCodeAgent`):

- A multi-agent workflow for language-agnostic translation and validation of large repositories.
- ARCMiS borrows: the `tool_projects` dataset and the four-stage pipeline discipline (ADR 0018, ADR 0022).

Claude Modernization Plugin:

- Discovery, brief, and execution flow with evidence-gated phase exits.
- ARCMiS borrows: the phase gate idea. A phase exit needs evidence, not a model claim.

System-one models: jev and laya

1. LLM-as-a-judge. A full model decode classifies an outcome. Cost: one model turn per judgment, on the same busy GPU.

2. Jev-as-a-judge. A local specialist checkpoint does the same judgment. In ARCMiS:

- `jev_judge.rs` : guard ask arbitration (ADR 0023).
- `jev_triage.rs` : per-dispatch triage of member results (ADR 0027).

3. Typed decisions (laya, arXiv:2609.26550). One forward pass answers `choice` , `score` , and `noul` questions. No token decoding. See `.omp/skills/laya` .

Bottleneck use - honest state: the harness wires per-dispatch triage with a min-confidence cascade (`TriageVerdict` from five `agent_trace_observability` questions). The round policy defaults to `observe` : verdicts get recorded, but no round stops on them. Enforcement (`policy: enforce`) exists in `round_triage/policy.rs` and is not the campaign default.

Benchmarks

Dataset and progress

Dataset: ReCodeAgent `tool_projects` - the campaign brief says 114 problems. Counted directories: **119** (alphatrans 4, crust 100, oxidizer 7, skel 8). Delta: 5 uncounted problems. Numbers below use the brief's 114.

Progress (live from `.artifacts/experiments/ROUNDS.yaml` and `SUMMARY.md`, read 2026-10-04):

- 44 ledger rounds: 19 SOLVED, 15 CLEARED, 15 NOT CLEARED, 10 ABORTED.
- Engine: q27 v0.14.1 Bonsai-2 T3-MTP-slim through the netmind provider path. Ninfer is retired.
- **13 of 114 problems cleared** (leftpad, amp, ulidgen, gonameparts, avalanche, colorsys, morton, libqueue, murmurhash, geofence, gfc, 2dpartint, bhshell). A problem counts once across retries.

Rate, estimate, and the ReCodeAgent record

Estimated time to complete (from the current rate; estimate, not a commitment):

- Campaign v3 started 2026-09-30. About 4.4 days elapsed to 2026-10-04.
- Rate: 13 cleared / 4.4 days = about 2.95 per day.
- Remaining: $(114 - 13) / 2.95$ = about 34 days.

Published ReCodeAgent record (different protocol):

Technique	Result	Protocol
ReCodeAgent	99.4% compile, 86.5% test pass	Claude 4.5 Sonnet, 118 projects, their developer-test suites
ARCMiS v3	13 of 114 problems cleared	Local qwen3.8-27B-class 27B engine, toolchain-only scoring, one set at a time

The two rows are directional context, not a like-for-like comparison. The model class, the budget, and the pass protocol differ.

References

- ReCodeAgent: A Multi-agent Workflow for Language-Agnostic Translation and Validation of Large-Scale Repositories. arXiv:2604.07341, ASE 2026.
- LLM-Based Multi-Agent Orchestration: A Survey of Frameworks, Communication Protocols, and Emerging Patterns. (the Q2 survey)
- Meta-Harness. arXiv:2603.28052.
- Typed decisions. arXiv:2609.26550. `laya-rs` crate, `.omp/skills/laya` .
- ARCMiS ADRs 0018, 0022, 0023, 0026, 0027, 0028 in `docs/adr/` .
- Dataset: `assets/ReCodeAgent/data/tool_projects` .
- Ledger: `.artifacts/experiments/ROUNDS.yaml` , `SUMMARY.md` .

